

February 1, 2026

Hiring Manager

Veeva Systems

Re: Senior Director of Engineering and Data Operations

I am writing to express my strong interest in the Senior Director of Engineering and Data Operations position for Veeva's OpenData team. With direct experience managing HCP and HCO healthcare reference data at Gilead Sciences, over 10 years architecting enterprise data platforms, and proven leadership of engineering and data operations teams in life sciences, I am excited about the opportunity to build and scale Veeva's healthcare reference data platform.

My experience is directly aligned with OpenData's mission. At Gilead Sciences, I led a team of ~5 people managing HCP and HCO healthcare reference data for Medical Affairs-healthcare professional and healthcare organization profiles were the primary dataset. I built Spotfire dashboards tracking HCP/HCO interactions, affiliations, engagements, and relationship networks, architected a centralized data lake consolidating HCP/HCO data from multiple sources, and managed vendors for data acquisition, cleansing, and standardization of HCP/HCO profiles. I established comprehensive data governance frameworks for healthcare reference data ensuring data quality, integrity, and compliance. This direct experience with HCP/HCO reference data-the same entity types OpenData manages-positions me uniquely to understand the data quality challenges, relationship complexities, and operational requirements of maintaining accurate healthcare professional and organization profiles at enterprise scale.

Beyond healthcare reference data, I have owned architecture, build, and delivery of large-scale data platforms across their full lifecycle. At Redica Systems, I architected a multi-tenant platform on Snowflake that acquires data from government sources (FDA inspection reports), implements data pipelines for ingestion and processing, performs curation and standardization, and enforces comprehensive data governance. At Takeda, I led engineering and data operations teams (3-7 people), architected Big Data solutions, and served as platform administrator for Spotfire, owning data modeling, ETL processes, enterprise integrations, and governance. My experience spans AWS, Snowflake, Databricks, and Azure, and I have deep expertise in data engineering, ETL/ELT pipelines, data warehousing, and building scalable cloud platforms that handle complex data source acquisition and quality management.

I am a hands-on technical leader who can help build and design platforms alongside teams. At Redica, I personally architected and deployed a production LLM classification system processing government data end-to-end. At Tidal Wave Analytics, I built production ML platforms on AWS Databricks and developed full-stack applications including backend data pipelines (Python), frontend applications (React/Node.js), and cloud data management (Supabase). I drive technical architecture decisions, translate complex requirements into sound technical design, and provide technical guidance and mentorship to foster cultures of technical excellence. I have experience hiring, developing, and leading high-performing teams in fast-paced life sciences environments (Regeneron, Takeda, Gilead, plus consulting for Eli Lilly and Merck). I collaborate effectively with Product Management, QA, and cross-functional R&D teams in Agile environments, and my life sciences background provides deep domain expertise in healthcare data operations and regulated environments. I would be excited to bring this combination of healthcare

reference data experience, data platform architecture expertise, and technical leadership to help Veeva build world-class HCP/HCO data products serving the life sciences industry.

I would welcome the opportunity to discuss how my experience managing HCP/HCO healthcare reference data and building enterprise data platforms can contribute to OpenData's growth. Thank you for considering my application.

Sincerely,

Matt Oremland, PhD